

To ensure your AI agent is correctly parsing the details from the **Aura Intelligence** knowledge base, you should test it with a mix of factual, troubleshooting, and "policy-boundary" questions.

Here are a few test cases categorized by their intent:

1. Factual & Pricing Questions

- "What is the main difference between Aura Pro and Aura Enterprise?"
- "I'm a freelancer. How many AI credits do I get per month with the \$29 plan?"
- "Does the company use my data to train its global AI models?"
- "I ran out of credits. Can I buy more without upgrading my whole plan?"

2. Technical & Integration Questions

- "How do I fix a sync issue between Aura and my Google Calendar?"
- "Will Aura tell people in a Zoom meeting that it is recording them?"
- "What is the API rate limit for an Enterprise user?"
- "Can I use Aura to automatically update my Salesforce leads?"

3. Policy & Security Questions

- "I work in a hospital. Are you HIPAA compliant?"
- "What happens if I use Aura to generate phishing emails?"
- "I want to delete all my data. How do I exercise my 'Right to be Forgotten'?"
- "If I cancel my subscription today, do I get a refund for the remaining two weeks of the month?"

4. Edge Cases & Logic (Advanced Testing)

- **The "Hallucination" Check:** "Can Aura give me legal advice that I can use in court without a lawyer?"
 - *Ideal Answer: No, the agent should reference the 'Human-in-the-Loop' policy on Page 7.*
- **The "Out of Credits" Math:** "I have 500 words of text and a 10-minute meeting to transcribe. How many credits will that cost me?"
 - *Ideal Answer: 15 credits (5 for text + 10 for transcription based on Page 3).*
- **Urgency Check:** "Our system is completely down and we are an Enterprise client. How fast will you respond?"
 - *Ideal Answer: 2 hours (based on the SLA on Page 8).*